

Mathematics of Deep Learning

Lecture 2: Expressive/Approximation Power of Neural Nets

Prof. Predrag R. Jelenković
Time: Wednesday 4:10-6:40pm

Dept. of Electrical Engineering
Columbia University , NY 10027, USA
Office: 812 Schapiro Research Bldg.
Phone: (212) 854-8174
Email: predrag@ee.columbia.edu
URL: <http://www.ee.columbia.edu/~predrag>

Last Lecture: General Learning Framework

- ▶ **Supervised learning:** there is an input-output relationship

$$Y = f(X)$$

- ▶ f - unknown
- ▶ Training data (observations): $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Objectives:
 - ▶ Learn/estimate \hat{f} from training data
 - ▶ Inference/prediction: Use \hat{f} to predict outcomes on unseen/new data
- ▶ Two problems:
 - ▶ Regression: Y is quantitative
 - ▶ Classification: Y is categorical
- ▶ **Unsupervised learning: Just X , no output Y (no labels)**
Typical problems:
 - ▶ Distribution estimation: learn the distribution/density, $p(x)$
 - ▶ Generative modeling: first estimate the distribution, then use it for other learning problems, e.g., QDA/LDA
 - ▶ Dimensionality reduction, e.g., PCA
 - ▶ Clustering and, in general, any data mining, e.g., data association, etc.

Last Lecture: Curse of Dimensionality

Why is it hard to estimate a function in high dimensions?

- ▶ How do we estimate a density of one dimensional X on $[0, 1]$?

Question: Say, we want to learn $p(x) : [0, 1] \rightarrow \mathbb{R}$

- ▶ Solution: split $[0, 1]$ in 100 bins of size $\epsilon = 0.01$, get about 1000 samples of X and plot a histogram

This should be a pretty good estimate of $p(x)$

- ▶ Suppose X is supported on 100 dimensional cube $[0, 1]^{100}$

Learn density $p(x) : [0, 1]^{100} \rightarrow \mathbb{R}$

- ▶ The preceding solution: splitting $[0, 1]^{100}$ in bins of size 0.01^{100} would require more than


10^{200} samples (!)

This is usually referred to as **curse of dimensionality**

- ▶ Many real world problems are high dimensional
Images $> 10^6$ dimensions; gene expression data $> 10,000$
- ▶ **Only hope:** existence of low dimensional structure

Overfitting Problem

One can fit infinitely many functions through a finite set of points, but



Overfitting:

- ▶ Low training error does not imply low testing error
- ▶ More complicated models not always better
 - ▶ Less interpretability
 - ▶ More difficult to train
- ▶ **Mystery:** Deep is learning highly flexible, possibly millions of parameters, but usually doesn't overfit.
- ▶ **John von Neumann elephant quote:** "With four parameters I can fit an elephant, and with five I can make him wiggle his trunk."

Last Lecture: General Supervised Learning Setup

Supervised learning problem can be formulated as (say, $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}^m$)

$$\hat{f}^* = \arg \min_{\hat{f} \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{f}(x_i)) + \lambda R(\hat{f}) \quad (1)$$

- ▶ $\ell : \mathbb{R}^{p+m} \rightarrow \mathbb{R}$ - loss function, and **empirical risk**/loss is defined as $\mathbf{x} = (x_1, \dots, x_p), \mathbf{y} = (y_1, \dots, y_m)$

$$\bar{L}(\mathbf{x}, \mathbf{y}) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{f}(x_i))$$

Hence, Equation (2) is called **Empirical Risk Minimization (ERM)**

- ▶ \mathcal{H} - Hypothesis class (class of approximation functions)
Desirable properties of \mathcal{H} :
 - ▶ Rich/versatile, yields accurate predictions, easy to train (e.g., problem (2) is convex), interpretable/simple, etc.
- ▶ $\lambda R(\hat{f}) \in \mathbb{R}^+$ - regularizer/penalty, shrinkage term
 - ▶ $\lambda R(\hat{f})$ - shrinking \mathcal{H} : if $\lambda_2 > \lambda_1 \rightarrow \mathcal{H}_{\lambda_2} \subset \mathcal{H}_{\lambda_1}$
 - ▶ Should prevent overfitting

Linear Examples: Regression and Classification

Regularized Linear Regression (say $y \in \mathbb{R}$)

- ▶ \mathcal{H} : set of all p -dimensional linear functions
 $\hat{f}(x) = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$
- ▶ Quadratic loss: $\ell(y_i, \hat{f}(x_i)) = (y_i - \hat{f}(x_i))^2$
- ▶ Typical penalties
 - ▶ Ridge - l_2 norm: $R(\hat{f}) = \langle \hat{f}, \hat{f} \rangle = \|\hat{f}\|^2 = \sum \beta_j^2$
 - ▶ LASSO: l_1 norm (basis pursuit): $R(\hat{f}) = \sum |\beta_j|$

Support Vector Classifier: Separate two classes by a hyperplane,
 $y_i \in \{-1, 1\}$

$$\min_{\beta_0, \beta} \sum_{i=1}^n \max[0, 1 - y_i(\beta_0 + x_{i1}\beta_1 + \cdots + x_{ip}\beta_p)] + \lambda \sum_{j=0}^p \beta_j^2$$

$\ell(x_i, y_i) = \max[0, 1 - y_i(\beta_0 + x_{i1}\beta_1 + \cdots + x_{ip}\beta_p)]$ is called **hinge loss**

- ▶ **What if the problem is not linear?**
This is most likely the case.


Nonlinear Functions: Transform/Expand Features

Example: Classifying red and blue point on a real line

- ▶ Can't be separated by a single point (hyperplane)



- ▶ Nonlinear feature transformation/expansion: $x \rightarrow (x, x^2)$
- ▶ Now, the points are separable by a single hyperplane (!)



- ▶ Feature engineering: $x \rightarrow \phi(x)$
- ▶ Well-developed theory: Reproducing Kernel Hilbert Spaces (RKHS)
- ▶ Problem: How do we find feature functions/kernels?

Where Does Deep Learning Fits in This Framework?

- ▶ Rich \mathcal{H} : Provides a versatile parametric class of functions
 - ▶ Universal function approximation: depth helps improves expressiveness
 - ▶ However, it is difficult to train: non-convex optimization
 - ▶ Often produces accurate predictions in practice, but difficult to understand and interpret
- ▶ Automatic feature extraction
 - ▶ Traditional Feature Engineering approach: expert constructs feature mapping $\phi : \mathcal{X} \rightarrow \Phi$. Then, apply machine learning to find a linear predictor on $\phi(\mathbf{x})$.
 - ▶ “Deep learning” approach: neurons in hidden layers can be thought of as features that are being learned automatically from the data
 - ▶ Shallow neurons corresponds to low level features, while deep neurons correspond to high level features

Parametric Supervised Learning

- ▶ \mathcal{H} is a parametric class of functions
 $f(w, x), w \in \mathcal{W}, w = (w_1, \dots, w_k)$
 - ▶ Examples: polynomials, or other linear combinations of basis, **neural networks**
- ▶ Finding $f \in \mathcal{H}$ is equivalent to finding $w \in \mathcal{W}$

Hence, our general supervised learning problem can be formulated in terms of w as (say, $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}^m, \ell : \mathbb{R}^{p+m} \rightarrow \mathbb{R}$ - loss function)

$$\hat{w} = \arg \min_{w \in \mathcal{W}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(w, x_i)) + \lambda R(w) \quad (2)$$

How do we solve the preceding problem?

- ▶ When the problem is convex and we are lucky, we can find an explicit \hat{w} by solving (e.g., generalized (Kernel) ridge regression)

$$\frac{\partial}{\partial w_i} \left(\frac{1}{n} \sum_{i=1}^n \ell(y_i, f(w, x_i)) + \lambda R(w) \right) = 0, \quad i = 1, \dots, k.$$

Parametric Supervised Learning

Let us denote the objective function

$$F(w) := \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(w, x_i)) + \lambda R(w)$$

In general, we use the following numerical algorithm:

1. Initialization: Pick an initial value w_0 , possibly random
2. Iteration: Follow the path of steepest descent = direction of negative gradient.

For the preceding procedure to find a local minimum

- ▶ Avoid getting stuck on a flat surfaces, say flat saddle points.
- ▶ If F is convex, we can find a global minimum.

Properties of gradient, $\nabla f(x)$:

- $\nabla f(x)$: direction of steepest ascent/max increase of $f(x)$
- $-\nabla f(x)$: direction of steepest descent/max decrease of $f(x)$
- $\nabla f(x)$ is perpendicular to level curves $f(x) = c$


Path of Gradient Descent

- $\mathbf{x}(t)$ - path of steepest gradient descent given initial condition $\mathbf{x}(0)$ is given by an ODE

$$\frac{d\mathbf{x}(t)}{dt} = -\eta \nabla f(\mathbf{x}(t)), \quad (3)$$

where η is the rate/speed of descent, a.k.a. learning rate. Note that $\mathbf{x}'(t)$ is tangent to the curve $\mathbf{x}(t)$, and thus parallel to $\nabla f(\mathbf{x}(t))$.

- Example: $f(\mathbf{x}) = ax_1^2 + bx_2^2, a, b > 0, \Rightarrow$
 $\nabla f(\mathbf{x}) = (2ax_1, 2bx_2) \Rightarrow \mathbf{x}'(t) = -2\eta a x_1(t), \mathbf{x}'_2(t) = -2\eta b x_2(t)$
 $x_1(t) = x_1(0)e^{-2\eta at}, \quad x_2(t) = x_2(0)e^{-2\eta bt}$



Gradient Descent Algorithm

GD Algorithm is a discrete linear approximation to Equation (3)

$$\frac{\mathbf{x}(t + \Delta t) - \mathbf{x}(t)}{\Delta t} \approx \frac{d\mathbf{x}(t)}{dt} = -\eta \nabla f(\mathbf{x}(t))$$

or equivalently (with a small abuse of notation)


$$\mathbf{x}^{(t+1)} = \mathbf{x}^{(t)} - \eta_t \nabla f(\mathbf{x}^{(t)})$$

- ▶ Hence, after initialization at $\mathbf{x}^{(0)}$, the GD Algorithm follows the preceding iteration to a local minimum
- ▶ Stopping criterion (could be): $|\mathbf{x}^{(t+1)} - \mathbf{x}^{(t)}| < \epsilon$

Adaptive GD (AdaGrad): modifies the learning rate η_t in each iteration t

More Interesting Landscape

- ▶ $f(x) = (x_1^2 - 1)^2 + (x_1^2 x_2 - x_1 - 1)^2$
- ▶ Gradient
$$\nabla f(x) = \begin{bmatrix} 4x_1(x_1^2 - 1) + 2(2x_1x_2 - 1)(x_1^2 x_2 - x_1 - 1) \\ 2x_1^2(x_1^2 x_2 - x_1 - 1) \end{bmatrix}$$
- ▶ Oscillations in "narrow valleys"



Motivation for **momentum**: remembers/averages previous Δx

$$\mathbf{x}^{(t+1)} = \mathbf{x}^{(t)} - \eta_t \nabla f(\mathbf{x}^{(t)}) + \mu_t (\mathbf{x}^{(t)} - \mathbf{x}^{(t-1)})$$

Stochastic Gradient Descent

- ▶ Stochastic approximation of gradient descent
- ▶ Function (typically encountered in learning)

$$f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$$

- ▶ Computationally expensive gradient for large n
- ▶ Approximation: pick a random subset $\mathcal{S} \in [1, n]$

$$\frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \nabla f_i(x)$$

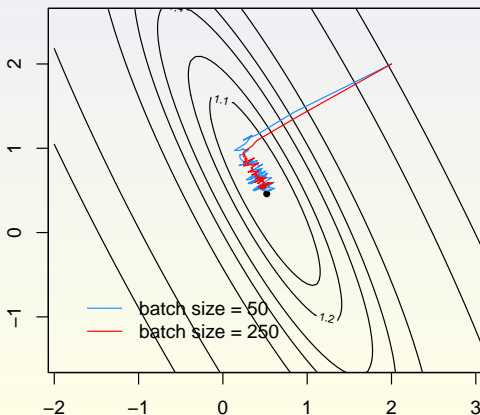
- ▶ \mathcal{S} - called batch/mini-batch
- ▶ Example
 - ▶ n scalar data points x_1, x_2, \dots, x_n
 - ▶ objective

$$\min_c \frac{1}{n} \sum_{i=1}^n (x_i - c)^2$$

SGD Example: Linear Regression

- ▶ Data: $(x_i, y_i)_{i=1}^n$, $n = 10^5$
- ▶ Loss function: $L(\beta, x, y) = \frac{1}{n} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$
- ▶ Stochastic gradient

$$\nabla_{\beta} \hat{L}(\beta, x, y) = \frac{2}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \begin{bmatrix} (\beta_0 + \beta_1 x_i - y_i) \\ x_i(\beta_0 + \beta_1 x_i - y_i) \end{bmatrix}$$




Math of Deep Learning: The Preliminary List of Topics

- ▶ **The preliminary list of topics** include the results on:
 - ▶ **Expressive (approximation) power of neural networks**
 - ▶ Depth separation results: Can deep neural networks of depth $(d + 1)$ express functions much more efficiently in terms of the number of neurons and parameters compared to networks of depth d ? What classes of functions can deep neural nets approximate well?
 - ▶ Connection between wide neural nets and Kernels: Neuro Tangent Kernels (NTK)
 - ▶ Global versus local optimality
 - ▶ Training and convergence properties of wide neural nets:
 - ▶ Is training converging to global min?
 - ▶ Is training converging far or close to the initial NTK?
 - ▶ Generalization error: Deep neural networks have a lot of parameters. How come they are not overfitting?
 - ▶ Basic generalization concepts from machine learning theory will be covered
 - ▶ Deep residual networks
 - ▶ Deep generative probabilistic models, e.g., deep Boltzmann machines (time permitting)
 - ▶ Etc.

Deep Neural Networks, a.k.a. Multilayer Perceptrons


- Feed-forward network



- Designed to mimic the function of [neurons](#)
- Blue nodes: activation functions/neurons
- Depth = lengths of a longest path
- Deep network: depth ≥ 3
- Very successful in solving practical problems

A Single Artificial Neuron

- ▶ A **single neuron** function: $\mathbf{x} \mapsto \sigma(b + \langle \mathbf{w}, \mathbf{x} \rangle)$, where $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is called the **activation function** of the neuron. Inner/dot product:
 $\langle \mathbf{w}, \mathbf{x} \rangle = \sum x_i w_i$.
- ▶ More compact notation $\langle \tilde{\mathbf{w}}, \tilde{\mathbf{x}} \rangle$, where $\tilde{\mathbf{x}} = (1, \mathbf{x})$, $\tilde{\mathbf{w}} = (b, \mathbf{w})$



- ▶ E.g., σ is a sigmoidal function



Common Activation Functions/Perceptrons

- ▶ step function

$$\sigma(x) = 1_{\{x>0\}} \quad \sigma'(x) = 0, x \neq 0$$

- ▶ logistic

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad \sigma'(x) = \sigma(x)(1 - \sigma(x))$$

- ▶ rectified linear unit (ReLU)


$$\sigma(x) = \max\{x, 0\} \quad \sigma'(x) = 1_{\{x>0\}}, x \neq 0$$

- ▶ soft-plus

$$\sigma(x) = \log(1 + e^x) \quad \sigma'(x) = \frac{1}{1 + e^{-x}}$$


Example

- ▶ Single neuron is a binary half-space classifier: $\text{sign}(w \cdot x + b)$
- ▶ 2 layer networks can express **intersection of halfspaces**



Example

- ▶ 3 layer networks can express **unions** of intersection of halfspaces



How do we train neural networks?

- ▶ Neural nets: excellent hypothesis class, but difficult to train
- ▶ Main technique: Stochastic Gradient Descent (SGD)
- ▶ Not convex, no guarantees, can take a long time, but:
 - ▶ Often still works fine, finds a good solution
 - ▶ With some luck:)

Stochastic Gradient Descent (SGD) for Neural Networks

Common Training Ideas:

- ▶ Random initialization: rule of thumb, $w[u \rightarrow v] \sim U[-c, c]$ where $c = \sqrt{3/|\{(u', v) \in E\}|}$ (or small Gaussian instead of $U[-c, c]$)
- ▶ Update step with Nesterov's momentum: Initialize $\theta = 0$ and:

$$\begin{aligned}\theta_{t+1} &= \mu_t \theta_t - \eta_t \tilde{\nabla} L(w_t + \mu_t \theta_t) \\ w_{t+1} &= w_t + \theta_{t+1}\end{aligned}$$

where:

μ_t is momentum parameter (e.g. $\mu_t = 0.9$ for all t)

η_t is learning rate (e.g. $\eta_t = 0.01$ for all t)

$\tilde{\nabla} L$ is an estimate of the gradient of L based on a small set of random examples (often called a “minibatch”)

- ▶ Efficient gradient calculation: [Backpropagation](#)

Expressive/Approximation Power of Neural Nets

Question: What types of functions can be approximated with neural networks?

Two approaches to answering this question:

- ▶ Constructive approach: for a given function, we construct explicitly its neural network approximation.
 - ▶ This is usually more intuitive, but less general.
- ▶ Existence approach: we prove that for a given class of functions, there exists an accurate neural network approximation.
 - ▶ This gives us confidence that good approximations exist, but less informative.

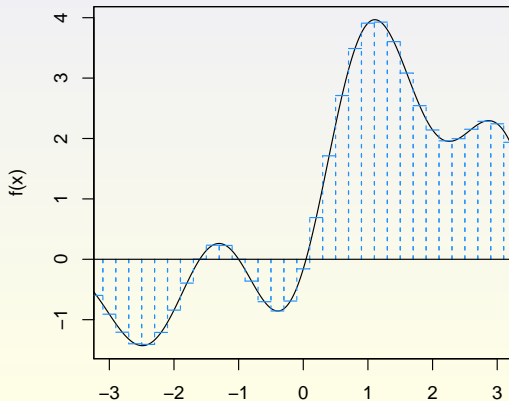
Expressive Power of Shallow Neural Nets

Questions: What type of real-valued functions can we approximate with NNs with one hidden layer? Such NN approximations can be written as:

$$\hat{f}(x) = \sum_{i=1}^n a_i \sigma(\langle w_i, x \rangle + b_i), \quad x, w \in \mathbb{R}^d, a_i, b_i \in \mathbb{R}$$

Single variable example: explicit approximation with simple/step unctnons

$$\sigma(x) = 1_{\{x>0\}}, \quad \phi(x) = 1_{\{x>0\}} - 1_{\{x>1\}} \quad (\text{simple/pulse function})$$




Expressive Power of Neural Nets

Piecewise Linear Approximation: ReLU - $\sigma(x) = \max(0, x)$

Example: $f(x) : [0, \pi] \rightarrow [0, 1]$

$$f(x) = \sin(x)$$

Piecewise Linear Approximation




Example: Function That Is Difficult to Approximate

Consider function $f(x) : [0, 1) \rightarrow [-1, 1]$

$$f(x) = \sin\left(\frac{1}{1-x}\right)$$

is continuous and infinitely differentiable.

Difficult to Approximate: Why?



Function Regularity: Uniform Continuity

So, we need to impose some regularity on $f(x)$.

Definition Let $I \subset \mathbb{R}$ be an interval. Function $f(x) : I \rightarrow \mathbb{R}$ is **uniformly continuous** if for any $\epsilon > 0$, there exists $\delta > 0$, such that

$$|x_1 - x_2| < \delta \quad \Rightarrow \quad |f(x_1) - f(x_2)| < \epsilon$$

Remarks

- ▶ If I is a closed interval, $[a, b]$, then

ordinary continuity \Leftrightarrow uniform continuity

- ▶ Hence, to prevent the problem from the preceding example, $f(x)$ needs to be defined on a **closed interval** $[0, 1]$ instead of $[0, 1)$.
- ▶ The preceding definition extends to \mathbb{R}^k spaces with Euclidean norm, and in general metric spaces.
- ▶ Finite closed intervals generalize to **compact sets**.


From Sigmoidal to Step Function

Starting with any sigmoidal functions, say

$$\sigma(x) = \frac{1}{1 + e^{-x}},$$

we can approximate arbitrarily close a step function by using $\sigma(wx)$ and w large enough.

$$\sigma(wx), w = 1, 3, 5, 8, 30$$



Making Simple/Pulses Function, a.k.a. Haar Wavelet

Haar scaling function (single pulse) is defined


$$\phi(x) = 1_{\{x>0\}} - 1_{\{x-1>0\}}$$

Two sigmoidal functions, we can create a perfect pulse (Haar function)

$$\phi(x) \approx \phi_w(x) = \sigma(wx) - \sigma(w(x-1))$$

for w large enough (which can be used to make Haar wavelet basis)

$$\phi_w(x), w = 3, 5, 8, 30, 100$$




Non-monotonic Sigmoidal to Step Function

Do we need sigmoidal functions to be monotonic? For example

$$\sigma(x) = \frac{1}{1 + e^{-x}} + \sin(4\pi x)e^{-|x|}$$

Again, we can approximate arbitrarily close a step function by using $\sigma(wx)$ and w large enough.

$$\sigma(wx), w = 1, 3, 5, 8, 30$$



Continuity can be relaxed to $\sigma(x)$ being bounded with limits at $\pm\infty$.

Universal Approximation Theorem in 1D

Definitions

- ▶ A function $\sigma(x) : \mathbb{R} \rightarrow \mathbb{R}$ is called **sigmoidal** if it is bounded and

$$\lim_{x \rightarrow -\infty} \sigma(x) = 0 \quad \text{and} \quad \lim_{x \rightarrow \infty} \sigma(x) = 1.$$

(Equivalence to step function: $\sigma(wx) \approx 1_{\{x>0\}}$ for large w .)

- ▶ $C([a, b])$: space of continuous functions, $f(x) : [a, b] \rightarrow \mathbb{R}$
without loss of generality, we consider $C([0, 1])$

Theorem (Constructive/Explicit Representation) Consider a sigmoidal function, σ and $f \in C([0, 1])$. For every $\epsilon > 0$, there exist an integer n and $w > 0$ (depending on n), such that for $x \in [0, 1]$

$$\hat{f}(x) \stackrel{\text{def}}{=} \sum_{k=1}^n (f(k/n) - f((k-1)/n)) \sigma(w(x - k/n)) + f(0) \sigma(w(x + 1/n)),$$

then

$$\sup_{0 \leq x \leq 1} |\hat{f}(x) - f(x)| < \epsilon.$$

Universal Approximation Theorem in 1D: Proof

- ▶ Let \hat{g} be a simple function approximation of f

$$\hat{g}(x) := \sum_{k=1}^n (f(k/n) - f((k-1)/n)) 1_{\{x-k/n > 0\}} + f(0) 1_{\{x+1/n > 0\}}$$

- ▶ Since $f(x)$ is continuous on $[0, 1]$, and thus uniformly continuous, we can choose n large enough, such that for any $x \in [0, 1]$

$$|f(x) - \hat{g}(x)| \leq \frac{\epsilon}{2}$$

- ▶ To complete the proof, we need to bound the error between $\hat{f}(x)$ and $\hat{g}(x)$, i.e.,

$$|f(x) - \hat{f}(x)| \leq |f(x) - \hat{f}(x)| + |\hat{f}(x) - \hat{g}(x)| = \frac{\epsilon}{2} + |\hat{f}(x) - \hat{g}(x)|$$

- ▶ To this end, let $M := \max(1 + \sup_x \sigma(x), \sup_x f(x))$

Universal Approximation Theorem in 1D: Proof

We have two parameters to choose, n, w ,

- ▶ n large enough to ensure

$$|f(k/n) - f((k-1)/n)| \leq \frac{\epsilon}{4M}$$

- ▶ w to make $\sigma(wx)$ arbitrarily close to a step function, i.e., we can choose w large enough such that for all $|x| \geq 1/n$

$$|\sigma(wx) - 1_{\{x>0\}}| \leq \frac{\epsilon}{4Mn}$$

and around $x = 0$, since $\sigma(x)$ is bounded, for $|x| \leq 1/n$

$$|\sigma(wx) - 1_{\{x>0\}}| \leq 1 + \sup_x \sigma(x) \leq M$$

Hence, for $x \in [(j-1)/n, j/n]$,


$$\begin{aligned} |\hat{f}(x) - \hat{g}(x)| &\leq \sum_{k \neq j} |f(k/n) - f((k-1)/n)| |\sigma(w(x - k/n)) - 1_{\{x - k/n > 0\}}| \\ &\quad + |f(j/n) - f((j-1)/n)| |\sigma(w(x - j/n)) - 1_{\{x - j/n > 0\}}| \\ &\quad + |f(0)| |\sigma(w(x + 1/n)) - 1_{\{x + 1/n > 0\}}| \\ &\leq \frac{\epsilon}{4M} (n-1) \frac{\epsilon}{4Mn} + \frac{\epsilon}{4M} M + \frac{|f(0)|\epsilon}{4Mn} < \frac{\epsilon}{2} \end{aligned}$$

Sigmoidal Approximation Example

Consider function $f(x) : [0, \pi] \rightarrow [0, 1]$

$$f(x) = \sin(x)$$

Step Functions



Expressiveness of ReLU in 1D: Also Universal

- ▶ Recall, ReLU activation: $\sigma(x) = \max\{0, x\} =: x^+$
- ▶ To show that ReLU is universal, we can define a sigmoid $\sigma_1(x) = x^+ - (x - 1)^+$ and use the preceding theorem
- ▶ Or we can do it directly using piece-wise linear approximation
- ▶ The following approximation can be made arbitrarily close to $f \in C([0, 1])$, i.e., $|f(x) - \hat{f}(x)| < \epsilon$, for large enough n

$$\hat{f}(x) = \sum_{k=0}^{n-1} w_k \sigma(x - k/n) + f(0),$$

where the slopes w_k are chosen such that $\hat{f}(k/n) = f(k/n)$, i.e.,

- ▶ $w_0 = n(f(1/n) - f(0))$
- ▶ w_1 is such that $f(2/n) = 2w_0/n + w_1/n + f(0)$, yielding

$$w_1 = n(f(2/n) - 2f(1/n) + f(0))$$

- ▶ And so on, we inductively select w_k to satisfy

$$f\left(\frac{k+1}{n}\right) = \hat{f}\left(\frac{k+1}{n}\right) = w_0 \frac{k+1}{n} + w_1 \frac{k}{n} + \dots + w_k \frac{1}{n} + f(0)$$


Functions Equivalent to ReLU

- ▶ Similarly to sigmoidal functions, through scaling, many other functions are equivalent to ReLU. For example, one can define ReLU-equivalent functions as: continuous with

$$\lim_{x \rightarrow -\infty} \sigma(x) = 0, \quad \lim_{x \rightarrow -\infty} \frac{\sigma(x)}{x} = 1.$$

- ▶ Then, as $w \rightarrow \infty$,
$$\frac{\sigma(wx)}{w} \rightarrow \max(0, x)$$

Scaling soft-plus $\sigma(x) = \log(1 + e^x)$ to ReLU




ReLU Approximation

Consider function $f(x) : [0, \pi] \rightarrow [0, 1]$

$$f(x) = \sin(x)$$

Piecewise Linear Functions



Historical Comments on ReLU

- ▶ $\max(0, x)$ has a long history in statistics
- ▶ It is called **linear spline basis** (or hinge function)
- ▶ Bias b : in $\max(0, x + b)$ is called **knot**
- ▶ One hidden layer NN with ReLU is a **free knot linear spline**
 - ▶ Studied for 50+ years
- ▶ In general, one considers **polynomial spline bases**

$$\max(0, x^k), \quad k \geq 0$$

e.g., see Chapter 5 of The Elements of Statistical Learning book by Hastie et al.

Constructive Extension to Multivariate Functions

- ▶ How about multivariate functions $f(x) : [0, 1]^d \rightarrow \mathbb{R}$?
- ▶ Can we find a **constructive** uniform approximation?
 - ▶ Use **simple functions** in many dimensions.
 - ▶ **Fourier series** to the rescue (idea from Chen, Chen&Liu 1992 [CCL92])

Consider periodic functions $f : \mathbb{R}^d \rightarrow \mathbb{R}$ with unit period in all directions, i.e.,

$$f(\mathbf{x} + \mathbf{m}) = f(\mathbf{x}), \quad \mathbf{x} \in \mathbb{R}^d, \mathbf{m} \in \mathbb{Z}^d$$

Equivalently, these functions can be considered on d -torus: $T^d = S \times \cdots S = (S)^d$, where S is a circle of circumference one.

- ▶ Note the functions on $[0, 1]^d$ can be extended to periodic functions on T^d .
- ▶ Let $C(T^d)$ and $C([0, 1]^d)$ be continuous real-valued functions on unit d -torus and d -cube, respectively.
- ▶ Note that $C(T^d) \subset C([0, 1]^d)$, but $C([0, 1]^d) \not\subset C(T^d)$ the inference $C([0, 1]^d) \subset C(T^d)$ was incorrectly used in [CCL92]

Finite Series Multidimensional Fourier Approximation

- Bochner-Riesz means, $R, \alpha \geq 0$,

$$\hat{f}_R(\mathbf{x}) := \sum_{\mathbf{m}: \|\mathbf{m}\|_2 < R, \mathbf{m} \in \mathbb{Z}^d} \left(1 - \frac{\|\mathbf{m}\|_2^2}{R^2}\right)^\alpha a_{\mathbf{m}} e^{2\pi i \mathbf{m} \cdot \mathbf{x}} \quad (4)$$

Then, the following uniform convergence theorem can be found in Chapter 7 of Stein&Weiss: see Corollary 2.15 on p. 256 and Theorem 2.11 (b) on p. 253.

Theorem If $f \in C(T^d)$ and $\alpha > (d-1)/2$, then $\hat{f}_R \rightarrow f$ **uniformly**, as $R \rightarrow \infty$, i.e., for any $\epsilon > 0$, there is $R > 0$, such that

$$\sup_{\mathbf{x} \in T^d} |f(\mathbf{x}) - \hat{f}_R(\mathbf{x})| < \epsilon.$$

- Note that when f is real valued, we can ignore the imaginary part in Equation (4), and write ($a_{\mathbf{m}} = a_{\mathbf{m}R} + i a_{\mathbf{m}I}$)

$$\hat{f}_R(\mathbf{x}) := \sum_{\mathbf{m}: \|\mathbf{m}\|_2 < R} \left(1 - \frac{\|\mathbf{m}\|_2^2}{R^2}\right)^\alpha (a_{\mathbf{m}R} \cos(2\pi i \mathbf{m} \cdot \mathbf{x}) - a_{\mathbf{m}I} \sin(2\pi i \mathbf{m} \cdot \mathbf{x})) \quad (5)$$

Fourier Series Recipe for NN Approximation

Consider continuous multivariable functions ($C(T^d)$). We can approximate these function using the following 2 step procedure:

1. Approximate uniformly with NNS 1D functions $\cos(u)$ and $\sin(u)$ for $|u| \leq 2\pi\sqrt{d}R$. (Note that $\mathbf{x} \in [0, 1]^d$ and $\|\mathbf{x}\|_1 \leq \sqrt{d}\|\mathbf{x}\|_2$ imply $|\mathbf{m} \cdot \mathbf{x}| \leq \sum |m_i| = \|\mathbf{m}\|_1 \leq \sqrt{d}\|\mathbf{m}\|_2 \leq \sqrt{d}R$.)
2. Replace the NN approximation of \sin / \cos into the Fourier approximation Equation (5)

Complexity: The preceding 2-step approximation has exponential complexity in d since the number of summands in Equation (5) is $\sum_{\|\mathbf{m}\|_2 < R} 1 \approx \text{Vol}(n\text{-Ball}) = O(R^d) = O(e^{d \log R})$
Curse of dimensionality

Uniform convergence: a brute-force way to ensure uniform convergence is to make the approximation of \sin / \cos in Step 1 above very precise with an error smaller than ϵ/R^d

More General Existence Results

- ▶ Consider multivariate functions $f(x) : [0, 1]^d \rightarrow \mathbb{R}$
(or some other compact domain)
- ▶ We will look at the results of Cybenko (1989) and Hornik (1991)
- ▶ These results are more general, but approximations are not explicitly constructed.
 - ▶ Instead, proofs by contradiction: elegant, but less intuitive
- ▶ Involve non-elementary mathematics: functional analysis

Results by Cybenko

Theorem (Cybenko (1989)) Let σ be a continuous sigmoidal function (recall $\lim_{x \rightarrow -\infty} \sigma(x) = 0$ and $\lim_{x \rightarrow \infty} \sigma(x) = 1$ and pick an $f \in C([0, 1]^d)$). Then, the set of approximation functions of the form

$$\hat{f}(x) := \sum_j \alpha_j \sigma(\langle \mathbf{w}_j, \mathbf{x} \rangle + b_j)$$

is dense in $C([0, 1]^d)$ with the metric $d(f, g) = \sup |f(x) - g(x)|$, $f, g \in C([0, 1]^d)$, i.e., for any $f \in C([0, 1]^d)$ and $\epsilon > 0$, there exists a NN approximation function $\hat{f}(x)$, such that

$$\sup_{x \in [0, 1]^d} |f(x) - \hat{f}(x)| < \epsilon.$$

Few Results From Functional Analysis

Hahn-Banach Extension Theorem If X is a normed vector space with linear subspace M and $x_0 \in X \setminus \overline{M}$, then there exists a continuous linear map $L : X \rightarrow \mathbb{R}$ with $L(x) = 0$ for all $x \in M$, $L(x_0) = 1$, and $\|L\| \leq d(M, x_0)$.

(E.g., see Theorem 3.5, p. 60 in Functional Analysis, 2nd ed., by Rudin.)

Use of this theorem in our context:

- ▶ Proof by contradiction: consider the subspace M given by $\{\hat{f}(x) = \sum \alpha_j \sigma(\langle w_j, x \rangle + b_j)\}$, and assume that its closure \overline{M} is not the entire space of functions $C([0, 1]^d)$.
- ▶ We conclude that there exists a continuous linear map L on our function space that restricts to 0 on \overline{M} but is not identically zero.
- ▶ Hence, to prove the desired result, it suffices to show that any continuous linear map L that is zero on M must be the zero map, implying contradiction.

Few Results From Functional Analysis

Riesz Representation Theorem A bounded linear functional $L : C(X) \rightarrow \mathbb{R}$ can be expressed as

$$L(f) = \int f d\mu(x),$$

where μ is a finite signed measure supported on X and $f \in C(X)$.

Simple motivating example in 1D: Consider functions of one variable, $C([0, 1])$, and $\sigma(x) = 1_{\{x>0\}}$.

Then, if the space spanned by \hat{f} , call it M , does not approximate all functions in $C([0, 1])$, then there exists a linear operator L on $C([0, 1])$, which is not identically zero.

But, by Riesz Theorem, for any $0 \leq a < b \leq 1$, and $f = 1_{\{x-a>0\}} - 1_{\{x-b>0\}} \in M$

$$L(f) = \int (1_{\{x-a>0\}} - 1_{\{x-b>0\}}) d\mu(x) = \int_a^b d\mu(x) = 0$$

Hence, $\mu \equiv 0$, which is a contradiction.

- The difficulty in the general proof is to extend this argument to many dimensions and general sigmoidal functions, σ .

Cybenko's Proof: Outline

Proof by contradiction:

- ▶ Consider the subspace M given by $\{\sum \alpha_j \sigma(\langle \mathbf{w}_j, \mathbf{x} \rangle + b_j)\}$
- ▶ Assume that its closure \overline{M} is not the entire space of functions $C([0, 1]^d)$. (If $\overline{M} = C([0, 1]^d)$, then we are done.)
- ▶ Hence, by Hahn-Banach Theorem, there exists a continuous linear map L on our function space $C([0, 1]^d)$ that restricts to 0 on \overline{M} but is not identically zero.
- ▶ Then, by Riesz Representation Theorem, this linear functional can be expressed as integral, for any $f \in C([0, 1]^d)$

$$L(f) = \int f d\mu(x)$$

- ▶ **Key difficulty** is to prove that, since $\sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) \in M$,

$$\int_{[0,1]^d} \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) d\mu(\mathbf{x}) = 0 \quad \Rightarrow \quad \mu \equiv 0,$$

which **implies** $L \equiv 0$ on entire $C([0, 1]^d)$, resulting in contradiction.

Lemma 1 in Cybenko

Lemma 1 Let μ be finite, signed measure on $[0, 1]^d$. For any bounded, sigmoidal function, σ , if, for all $\mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}$

$$\int_{[0,1]^n} \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) d\mu(\mathbf{x}) = 0 \quad \Rightarrow \quad \mu \equiv 0$$

Proof Recall the step function approximation: as $\lambda \rightarrow \infty$, and then $\phi \rightarrow \infty$

$$\begin{aligned} \sigma(\lambda(\langle \mathbf{w}, \mathbf{x} \rangle + b) + \phi) &\rightarrow 1_{\{\langle \mathbf{w}, \mathbf{x} \rangle + b > 0\}} + \sigma(\phi) 1_{\{\langle \mathbf{w}, \mathbf{x} \rangle + b = 0\}} \\ &\rightarrow 1_{\{\langle \mathbf{w}, \mathbf{x} \rangle + b \geq 0\}} \end{aligned}$$

Hence, by Dominated Convergence Theorem, as $\lambda \rightarrow \infty$, and then $\phi \rightarrow \infty$

$$\begin{aligned} 0 &= \int_{[0,1]^d} \sigma(\lambda(\langle \mathbf{w}, \mathbf{x} \rangle + b)) d\mu(\mathbf{x}) = \int_{[0,1]^n} 1_{\{\langle \mathbf{w}, \mathbf{x} \rangle + b > 0\}} d\mu(\mathbf{x}) \\ &= \mu(\{\mathbf{x} : \langle \mathbf{w}, \mathbf{x} \rangle + b \geq 0\}), \end{aligned}$$

i.e., μ is zero on all half-spaces. (If μ were positive, we would be done.)

Lemma 1 in Cybenko

Proof - continued Since indicator (simple) functions are dense in $L^\infty(\mathbb{R})$ (bounded functions: $\mathbb{R} \rightarrow \mathbb{R}$), we conclude that the linear operator

$$F(h) = \int_{[0,1]^d} h(\langle \mathbf{w}, \mathbf{x} \rangle) d\mu(\mathbf{x}) \equiv 0$$

for any $h \in L^\infty(\mathbb{R})$

Hence, by choosing h to be $\sin(\langle \mathbf{w}, \mathbf{x} \rangle)$ and $\cos(\langle \mathbf{w}, \mathbf{x} \rangle)$, we get

$$\begin{aligned} & \int_{[0,1]^d} (\cos(\langle \mathbf{w}, \mathbf{x} \rangle) + i \sin(\langle \mathbf{w}, \mathbf{x} \rangle)) d\mu(\mathbf{x}) \\ &= \int_{[0,1]^d} \exp(i \langle \mathbf{w}, \mathbf{x} \rangle) d\mu(\mathbf{x}) = 0 \end{aligned}$$

for all \mathbf{w} .

Finally, Fourier transform of μ being zero, implies

$$\mu \equiv 0.$$

This completes the proof of the lemma, and, in view of the prior outline, the proof of Theorem 2 in Cybenko.

Hornik's Extensions

Hornik follows the general plan in Cybenko, which we outlined earlier. The key technical result is the generalization of the preceding Lemma 1 of Cybenko to allowing the activation functions, $\sigma(x)$, to be **bounded and nonconstant**, not necessarily sigmoidal.

Theorem 5 (Hornik) Let μ be finite, signed measure on $[0, 1]^d$. For any **bounded and nonconstant** σ , if, for all $\mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}$

$$\int_{[0,1]^n} \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) d\mu(\mathbf{x}) = 0 \quad \Rightarrow \quad \mu \equiv 0$$

Proof (sketch): As in Cybenko, the key idea is to show that the Fourier transform of μ is identically zero, which implies $\mu \equiv 0$.

We want to show that, for any \mathbf{w}, b

$$\int_{\mathbb{R}^n} \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b) d\mu(\mathbf{x}) = 0 \quad \Rightarrow \quad \mu \equiv 0.$$

First, we reduce the integration from \mathbb{R}^d to \mathbb{R} , by defining measures $\mu_{\mathbf{w}}$ on $\mathbb{R}, B \subset \mathbb{R}$, as

$$\mu_{\mathbf{w}}(B) = \mu(\mathbf{x} : \langle \mathbf{w}, \mathbf{x} \rangle \in B)$$

Hornik's Extensions

and observe that the prior integral becomes

$$\int_{\mathbb{R}^d} \sigma(\lambda \langle \mathbf{w}, \mathbf{x} \rangle + b) d\mu(\mathbf{x}) = \int_{\mathbb{R}} \sigma(\lambda t + b) d\mu_{\mathbf{w}}(t) = 0$$

Moreover, if we can show that $\mu_{\mathbf{w}} \equiv 0$ for each \mathbf{w} , then $\mu \equiv 0$ (“a measure is defined by all of its projections”), since then

$$\hat{\mu}(\mathbf{w}) = \int_{\mathbb{R}^d} \exp(i \langle \mathbf{w}, \mathbf{x} \rangle) d\mu(\mathbf{x}) = \int_{\mathbb{R}} \exp(it) d\mu_{\mathbf{w}}(t) = 0 \Rightarrow \hat{\mu} = 0 \Rightarrow \mu = 0.$$

(Note that we used the finiteness of μ here.)

Hence, the goal is to find arguments that justify the preceding equation.

Hornik's Extensions

To this end, use the convolution trick (that also uses the finiteness of μ). By convolving $\mu_{\mathbf{w}}(t)$ with a Gaussian $v(t) = e^{-t^2}$, we obtain a measure that has a density, letting us work with Lebesgue measure. $\mu_{\mathbf{w}} * v$ is also finite.

Next,

$$\begin{aligned} 0 &= \int_{\mathbb{R}} \left[\int_{\mathbb{R}} \sigma(\lambda t + (b + \lambda s)) d\mu_{\mathbf{w}}(t) \right] v(s) ds \\ &= \int_{\mathbb{R}} \int_{\mathbb{R}} \sigma(\lambda(t + s) + b) d\mu_{\mathbf{w}}(t) v(s) ds d\mu_{\mathbf{w}}(t) \\ &= \int_{\mathbb{R}} \sigma(\lambda t + b) d(v * \mu_{\mathbf{w}})(t); \end{aligned}$$

let $h(t) = (v * \mu_{\mathbf{w}})(t)$ is the convolution of $\mu_{\mathbf{w}}(t)$ and $v(t)$.

Then, using the results from abstract Fourier analysis, one shows that $h \equiv 0$. (see: Fourier Analysis on Groups, Rudin 1967)

This implies that the Fourier transform $\hat{h} = \hat{v} \hat{\mu}_{\mathbf{w}} = 0$.

Since, \hat{v} has no zeros, this results in $\hat{\mu}_{\mathbf{w}} \equiv 0 \Rightarrow \hat{\mu} \equiv 0 \Rightarrow \mu \equiv 0$, concluding the proof of Theorem 5 in Hornik.

Hornik's Extensions

Using Theorem 5 with more general activation functions, Hornik obtains the following results:

- ▶ **Theorem 1** If σ is unbounded and nonconstant, then NN approximations are dense in $L^p(\mu)$ for all finite measures μ on \mathbb{R}^d . $L^p(\mu)$, $p \geq 1$, is the space of functions f with $\int |f|^p d\mu < \infty$ and distance metric $d(f, g) = (\int |f - g|^p d\mu)^{1/p}$.
- ▶ **Theorem 2** If σ is continuous, bounded and nonconstant, then NN approximation are dense in the space of all continuous functions, $C(X)$, with compact domain $X \subset \mathbb{R}^k$ and supremum distance $d(f, g) = \sup_{x \in X} |f(x) - g(x)|$. (This is a full generalization of Theorem 2 in Cybenko.)
- ▶ **Theorems 3&4** Extend results to Sobolev spaces under the ℓ_p , $1 \leq p < \infty$ and supremum norm.

Sobolev spaces contain function that have up to m derivatives and distance is measured between functions and their derivatives.

Reading on the Universal Approximation Results

Shallow (1 hidden layer) networks are universal approximators

- ▶ Constructive proofs:
 - ▶ [A Constructive Proof and An Extension of Cybenko's Approximation Theorem](#), by Chen, Chen & Ruey-wen Liu, 1992. (Note: Not more general than Cybenko: requires continuity of periodic function extensions.)
 - ▶ [Constructive Approximation by Superposition of Sigmoidal Functions](#), by Costarelli and Spigler, 2013. (1D&2D)
- ▶ General proofs:
 - ▶ [Approximations by superpositions of sigmoidal functions](#), by Cybenko, 1989.
 - ▶ [Approximation capabilities of multilayer feedforward networks](#), by Hornik, 1991.

Reading For Next Class

Bounds on NN approximations:

- ▶ [Universal approximation bounds for superpositions of a sigmoidal function](#), by Barron 1993.
- ▶ [Approximation theory of the mlp model in neural networks](#), by Pinkus, 1999.

Expressive Power Depth Depth separation results: Can deep neural networks of depth $(d + 1)$ express functions much more efficiently in terms of the number of neurons compared to networks of depth d ?

- ▶ [Representation Benefits of Deep Feedforward Networks](#), by Telgarsky, 2015.
- ▶ [Error bounds for approximations with deep ReLU networks](#), by Yarotsky, 2017.
- ▶ [WHY DEEP NEURAL NETWORKS FOR FUNCTION APPROXIMATION?](#), Lianf and Srikant, 2017.
- ▶ [NEW ERROR BOUNDS FOR DEEP RELU NETWORKS USING SPARSE GRIDS](#), by Montanelli and Du, 2018.
- ▶ [The Power of Depth for Feedforward Neural Networks](#), by Eldan and Shamir, 2016.

General references on deep learning:

[UML] Chapter 14: Stochastic Gradient Descent
Chapter 20: Neural Networks

[DL] Chapter 6: Deep Feedforward Networks

Software: Download R, R Studio and TensorFlow
Or, equivalent for Python

If you would like to refresh or learn the concepts from real analysis, you could check Chapter 2 on basic topology concepts and Chapter 4 for continuity/uniform continuity in

- ▶ Principles of Mathematical Analysis, by W. Rudin.

More advanced reading on Fourier analysis

- ▶ Introduction to Fourier Analysis on Euclidian Spaces, by Stein and Weiss, 1971. (Chapter 7 is on multivariate Fourier series and Bochner-Riesz means.)
- ▶ Fourier Analysis on Groups, by W. Rudin. (Chapter 7 is used in Hornik.)