

Mathematics of Deep Learning

Lecture 12: VC Dimension and Generalization, Implicit Bias of Gradient Descent, ResNets and ODEs

Prof. Predrag R. Jelenković
Time: Tuesday 4:10-6:40pm

Dept. of Electrical Engineering
Columbia University , NY 10027, USA
Office: 812 Schapiro Research Bldg.
Phone: (212) 854-8174
Email: predrag@ee.columbia.edu
URL: <http://www.ee.columbia.edu/~predrag>

Learning Infinite Hypothesis Classes

For most hypothesis classes $|\mathcal{H}| = \infty$: What can be done here?

Common measures of complexity of hypothesis classes

- ▶ **Vapnik-Chervonenkis (VC) dimension**
Chapter 6, 28 & 20 in Shai's [ML2014] book
- ▶ **Rademacher Complexity**
Chapter 26 in [ML2014] book; this was recently used in
 - ▶ **A Priori Estimates For Two-layer Neural Networks**, by Weinan et al., Jan 2019.
 - ▶ **Fine-Grained Analysis of Optimization and Generalization for Overparameterized Two-Layer Neural Networks**, by Arora et al., Jan 2019.
- ▶ **PAC - Bayes**: Chapter 31 in [ML20014]
used recently in several NN papers, e.g.: see Neyshabur et al.
- ▶ **Compression Bounds**: Chapter 30 in [ML2014]

VC Dimension

- ▶ **VC dimension**: common characterization of sample complexity
- ▶ Introduced by Vapnik & Chervonenkis (VC)
- ▶ Can be used to characterize the sample complexity of NNs
- ▶ This is an example of bounding the generalization by \hat{f} instead of the target function f^* .

VC Dimension: Definition

- ▶ Let $C = \{x_1, \dots, x_{|C|}\} \subset \mathcal{X}$
- ▶ Let \mathcal{H}_C be the restriction of \mathcal{H} to C , namely,
 $\mathcal{H}_C = \{h_C : h \in \mathcal{H}\}$ where $h_C : C \rightarrow \{0, 1\}$ is s.t.
 $h_C(x_i) = h(x_i)$ for every $x_i \in C$
- ▶ Observe: we can represent each h_C as the vector
 $(h(x_1), \dots, h(x_{|C|})) \in \{\pm 1\}^{|C|}$
- ▶ Therefore: $|\mathcal{H}_C| \leq 2^{|C|}$
- ▶ We say that \mathcal{H} shatters C if $|\mathcal{H}_C| = 2^{|C|}$
- ▶ $\text{VCdim}(\mathcal{H}) = \sup\{|C| : \mathcal{H} \text{ shatters } C\}$
- ▶ That is, the VC dimension is the maximal size of a set C such that \mathcal{H} gives no prior knowledge w.r.t. C

VC dimension — Examples

To show that $\text{VCdim}(\mathcal{H}) = d$ we need to show that:

1. There exists a set C of size d which is shattered by \mathcal{H} .
2. Every set C of size $d + 1$ is not shattered by \mathcal{H} .

VC dimension — Examples

Threshold functions: $\mathcal{X} = \mathbb{R}$, $\mathcal{H} = \{x \mapsto \text{sign}(x - \theta) : \theta \in \mathbb{R}\}$

- ▶ Show that $\{0\}$ (or any other one-point set) is shattered
- ▶ Show that any two points cannot be shattered

VC dimension — Examples

Intervals: $\mathcal{X} = \mathbb{R}$, $\mathcal{H} = \{h_{a,b} : a < b \in \mathbb{R}\}$, where $h_{a,b}(x) = 1$ iff $x \in [a, b]$

- ▶ Show that $\{0, 1\}$ is shattered
- ▶ Show that any three points cannot be shattered
- ▶ Note that \mathcal{H} is a 2-parameter class and $\text{VCdim}(\mathcal{H}) = 2$

VC dimension — Examples


Axis aligned rectangles: $\mathcal{X} = \mathbb{R}^2$,

$\mathcal{H} = \{h_{(a_1, a_2, b_1, b_2)} : a_1 < a_2 \text{ and } b_1 < b_2\}$, where


$h_{(a_1, a_2, b_1, b_2)}(x_1, x_2) = 1$ iff $x_1 \in [a_1, a_2]$ and $x_2 \in [b_1, b_2]$

Show:

Shattered



Not Shattered



Note that \mathcal{H} is a 4-parameter class and $\text{VCdim}(\mathcal{H}) = 4$

VC dimension — Examples

Finite classes:

- ▶ Show that the VC dimension of a finite \mathcal{H} is at most $\text{VCdim}(\mathcal{H}) \leq \log_2(|\mathcal{H}|)$
 - ▶ since C cannot be shattered if $|\mathcal{H}| < 2^{|C|}$
- ▶ There can be arbitrary gap between $\text{VCdim}(\mathcal{H})$ and $\log_2(|\mathcal{H}|)$
 - ▶ e.g., consider $\mathcal{X} = \{1, 2, \dots, k\}$ and consider $\mathcal{H} = \{\text{step functions on } \mathcal{X}\}$
 - ▶ Then, $|\mathcal{H}| = k$, but $\text{VCdim}(\mathcal{H}) = 1$

VC dimension — Examples

Halfspaces: $\mathcal{X} = \mathbb{R}^d$, $\mathcal{H} = \{\mathbf{x} \mapsto \text{sign}(\langle \mathbf{w}, \mathbf{x} \rangle) : \mathbf{w} \in \mathbb{R}^d\}$

- ▶ Show that $\{\mathbf{e}_1, \dots, \mathbf{e}_d\}$ is shattered
- ▶ Show that any $d + 1$ points cannot be shattered
- ▶ Hence, $\text{VCdim}(\mathcal{H}) = d$

- ▶ Note again that \mathcal{H} is a d -parameter class and $\text{VCdim}(\mathcal{H}) = d$
- ▶ In general, one can expect that the VC dimension of a hypothesis class is equal to the number of parameters.

The Fundamental Theorem of Statistical Learning

Theorem (Theorem 6.8 in [ML] book)

Let \mathcal{H} be a hypothesis class of binary classifiers with

$\text{VCdim}(\mathcal{H}) = d$. Then, there are absolute constants C_1, C_2 , s.t.

1. \mathcal{H} is (agnostic) PAC learnable with sample complexity

$$C_1 \frac{d + \log(1/\delta)}{\epsilon^2} \leq m_{\mathcal{H}}(\epsilon, \delta) \leq C_2 \frac{d + \log(1/\delta)}{\epsilon^2}$$

2. (Realizable case) \mathcal{H} is PAC learnable with sample complexity

$$C_1 \frac{d + \log(1/\delta)}{\epsilon} \leq m_{\mathcal{H}}(\epsilon, \delta) \leq C_2 \frac{d \log(1/\epsilon) + \log(1/\delta)}{\epsilon}$$

Furthermore, this sample complexity is achieved by the ERM rule.

- ▶ We gave a sketch of the proof of part 2: the realizable case.
- ▶ **Part 1. will be covered today.** It is based on Rademacher complexity: more specifically Massarat Lemma 26.8 (see Chapter 28 of the book).

Sauer-Shelah-Perles Lemma

Let

$$\tau_{\mathcal{H}}(m) := \max_{C \in \mathcal{X}: |C|=m} |\mathcal{H}_C|$$

Lemma (Sauer-Shelah-Perles)

Let \mathcal{H} be a hypothesis class with $\text{VCdim}(\mathcal{H}) \leq d < \infty$. Then, for all $C \subset \mathcal{X}$ s.t. $|C| = m > d + 1$ we have

$$\tau_{\mathcal{H}}(m) \leq \sum_{i=0}^d \binom{m}{i} \leq \left(\frac{em}{d}\right)^d$$

- ▶ The lemma shows that the maximum number of different vectors that our hypothesis class \mathcal{H} can generate on a data set with $m > d + 1$ points grows polynomially, rather than exponentially in m .
- ▶ The proof is by induction in m , see p. 49 in the [ML] book.
- ▶ A more intuitive proof can be found in [Bartlett](#)

Proof of the Upper Bound of the Agnostic Case

- ▶ Theorem 6.8-1
- ▶ By Sauer's lemma, if $\text{VCdim}(\mathcal{H}) = d$, then

$$|\{(h(x_1), \dots, h(x_m)) : h \in \mathcal{H}\}| \leq \left(\frac{em}{d}\right)^d$$

- ▶ Denote

$$A = \{(1_{\{h(x_1) \neq y_1\}}, \dots, 1_{\{h(x_m) \neq y_m\}}) : h \in \mathcal{H}\}$$

and observe

$$|A| \leq \left(\frac{em}{d}\right)^d$$

Proof of the Upper Bound of the Agnostic Case

- Now, use the preceding estimate and recall the Massarat Lemma 26.8

$$\begin{aligned} R(A) &\leq \max_{\mathbf{a} \in A} \|\mathbf{a}\|_2 \frac{\sqrt{2 \log(|A|)}}{m} \\ &\leq \sqrt{m} \frac{\sqrt{2d \log(em/d)}}{m} = \sqrt{\frac{2d \log(em/d)}{m}} \end{aligned}$$

since $\|\mathbf{a}\|_2 \leq \sqrt{m}$

- Finally, the preceding bound and Theorem 26.5 yield, for any $h \in \mathcal{H}$ with probability at least $1 - \delta$

$$L(h) - \hat{L}_S(h) \leq \sqrt{\frac{8d \log(em/d)}{m}} + \sqrt{\frac{2 \log(1/\delta)}{m}}$$

VC Dimension of Neural Networks

Recall the graph notation for NNs:

- ▶ A neural network is obtained by connecting many neurons together
- ▶ We focus on feedforward networks, formally defined by a directed acyclic graph $G = (V, E)$
- ▶ Input nodes: nodes with no incoming edges
- ▶ Output nodes: nodes without outgoing edges
- ▶ weights: $w : E \rightarrow \mathbb{R}$
- ▶ Calculation using breadth-first-search (BFS), where each neuron (node) receives as input:


$$a[v] = \sum_{u \rightarrow v \in E} w[u \rightarrow v] o[u]$$

and output

$$o[v] = \sigma(a[v])$$

Multilayer Neural Networks

- ▶ Neurons are organized in layers: $V = \cup_{t=0}^T V_t$, and edges are only between adjacent layers
- ▶ Example of a multilayer neural network of depth 3 and size 6



Neural Network Hypothesis Class

- ▶ Given a neural network (V, E, σ, w) , we obtain a hypothesis $h_{V,E,\sigma,w} : \mathbb{R}^{|V_0|-1} \rightarrow \mathbb{R}^{|V_T|}$
- ▶ We refer to (V, E, σ) as the **architecture**, and it defines a hypothesis class by

$$\mathcal{H}_{V,E,\sigma} = \{h_{V,E,\sigma,w} : w \text{ is a mapping from } E \text{ to } \mathbb{R}\} .$$

- ▶ The architecture is our “Prior knowledge” and the learning task is to find the weight function w

Neural Network Sample Complexity

- ▶ **Theorem 1:** (σ - step/sign function) The VC dimension of $\mathcal{H}_{V,E,\text{sign}}$ is $O(|E| \log(|E|))$.
- ▶ **Theorem 2:** (σ - any sigmoidal function) The VC dimension of $\mathcal{H}_{V,E,\sigma}$, for σ being the sigmoidal function, is $\Omega(|E|^2)$.
- ▶ **Representation trick:** In practice, we only care about networks where each weight is represented using $O(1)$ bits, and therefore the VC dimension of such networks is $O(|E|)$, no matter what σ is.

Neural Network Sample Complexity

Proof of Theorem 1:

- ▶ Let $\tau_{\mathcal{H}}(m) = \max_{C \in \mathcal{X}: |C|=m} |\mathcal{H}_C|$, where \mathcal{H}_C is the restriction to C of binary valued functions in \mathcal{H}
- ▶ NN has T layers: $0, 1, 2, \dots, T$ with V_t nodes at layer t .
- ▶ Then, \mathcal{H} can be written as a composition

$$\mathcal{H} = \mathcal{H}^{(T)} \circ \dots \circ \mathcal{H}^{(1)}$$

- ▶ Furthermore, each class $\mathcal{H}^{(t)}$ can be decomposed per each neuron

$$\mathcal{H}^{(t)} = \mathcal{H}^{(t,1)} \times \dots \times \mathcal{H}^{t,|V_t|}$$

Neural Network Sample Complexity

Proof of Theorem 1:

- ▶ Then

$$\tau_{\mathcal{H}^{(t)}}(m) \leq \prod_{i=1}^{|V_t|} \tau_{\mathcal{H}^{(t,i)}}(m)$$

- ▶ Let $d_{t,i}$ be the number of edges that are headed to the i th neuron of layer t .
- ▶ Since each neuron is a homogenous half-space hypothesis class and the VC dimension of the the half-spaces is the dimension of their input, by Sauer's lemma,

$$\tau_{\mathcal{H}^{(t,i)}}(m) \leq \left(\frac{em}{d_{t,i}} \right)^{d_{t,i}} \leq (em)^{d_{t,i}}$$

implying

$$\tau_{\mathcal{H}}(m) \leq (em)^{\sum_{i,t} d_{t,i}} = (em)^{|E|}$$

Neural Network Sample Complexity

Proof of Theorem 1:

- Now, if we assume that m points are shattered, we must have

$$2^m \leq (em)^{|E|}$$

implying

$$m \leq |E| \log(em) / \log(2)$$

resulting in

$$m \leq O(|E| \log(|E|)),$$

which concludes the proof.

Generalization Bound for Unregularized NNs

Theorem Let $\mathcal{H} = (V, E, \sigma)$ be a hypothesis class of binary classifiers of multilayer NN with step function activation σ . Then, there are absolute constants C_1, C_2 , s.t. \mathcal{H} is (agnostic) PAC learnable with sample complexity

$$C_1 \frac{|E| \log(|E|) + \log(1/\delta)}{\epsilon^2} \leq m_{\mathcal{H}}(\epsilon, \delta) \leq C_2 \frac{|E| \log(|E|) + \log(1/\delta)}{\epsilon^2}$$

Furthermore, this sample complexity is achieved by the ERM rule.

- ▶ If σ is any sigmoid, $|E| \log(|E|)$ should be replaced by $|E|^2$
- ▶ Hence, we need **either regularization/shrinkage or prior knowledge on the target function** to reduce the sample complexity: e.g., Weinan et. al (2019), Neyshabur et. al. (2015-)

Recent Results on VC-dim of NNs With ReLUs

- ▶ Nearly-tight VC-dimension and Pseudodimension Bounds for Piecewise Linear Neural Networks, Bartlett et al., 2019.
- ▶ Compute tight upper and lower bounds on the VC-dimension of deep neural networks with the ReLU activation function.
- ▶ Let W be the number of weights and L be the number of layers. Then, the paper
 - ▶ proves that the VC-dimension is $O(WL \log(W/L))$
 - ▶ provides examples with VC-dimension $\Omega(WL \log(W/L))$
- ▶ Roughly $\text{VCdim}(\mathcal{H}) \approx WL$

More Results on GD Convergence and Generalization

A Comparative Analysis of the Optimization and Generalization Property of Two-layer Neural Network and Random Feature Models Under GD,
Weinan et al., Feb, 2020.

Key new results:

- Considers both lazy ($\beta = 1/\sqrt{m}$) and active ($\beta = 1/m$) training.

Theorem 3.2: Relaxed the over-parametrization assumption $m = \Omega(n^6/(\delta^3 \lambda_0^4))$ of Du et al. (2018) to $m = \Omega(n^2/(\delta \lambda_0^4))$

Theorem 3.3: Show that in the over-parametrized regime, $m = \Omega(n^2/(\delta \lambda_0^4))$, the approximations of the network where w -s are trained and the one with random initial w_0 -s are arbitrarily close, i.e.: **training yields the same approximation as a fixed kernel method.**

Theorem 4.1: ($\beta = c/m$) Relax the over-parametrization, but assume Barron function, and show that early stopping solution, Corollary 4.3, can be close to optimal for mildly over-parametrised networks.

Recall the Notation

- ▶ Training set: $S = \{(x_i, y_i)\}_{i=1}^n$; i.i.d. from a distribution $\rho_{x,y}$
- ▶ True (target) function: $f^*(x) = \mathbb{E}[y|x]$, where $y = f(x) + \xi$
- ▶ $f^*(x) : [-1, 1]^d \rightarrow [0, 1]$
- ▶ Two-layer neural network

$$f(x; \theta) = \sum_{k=1}^m a_k \sigma(w_k^\top x)$$

where $w_k \in \mathbb{R}^d$, $a_k \in \mathbb{R}$ and $\theta = \{(a_k, w_k)\}_{k=1}^m$

- ▶ **Scaling:** $\mathbb{P}[a_k(0) = \pm\beta] = 1/2$
 - ▶ β can depend on m , e.g.: $\beta = 1/\sqrt{m}$, or $\beta = 1/m$.
- ▶ $\sigma(x) : \mathbb{R} \rightarrow \mathbb{R}$: activation function
 $\sigma(x)$ scale free: $\sigma(\alpha x) = \alpha \sigma(x)$, $\alpha \geq 0$, $x \in \mathbb{R}$
e.g., ReLU or Leaky ReLU

Training

- ▶ Loss function: $\ell(y, y') = (y - y')^2/2$
- ▶ Ultimate goal: minimize the **population (true) risk**

$$L(\theta) = \mathbb{E}_{x,y}[\ell(f(x; \theta), y)]$$

- ▶ In practice: minimize the **empirical risk**

$$\hat{L}_n(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i; \theta), y_i)$$

- ▶ **Barron function**: A function $f : \Omega \rightarrow \mathbb{R}$ is called a Barron if it admits the following representation

$$f(x) = \int_{S^d} a(w) \sigma(w^\top x) d\pi(w),$$

where π is a probability distribution over $S^d = \{x : \|x\|_1 = 1\}$ and $a(\cdot)$ is a scalar function.

Barron Space

- ▶ **Barron norm:** Let f be a Barron function.

- ▶ Denote by Θ_f all possible representations of f

$$\Theta_f = \left\{ (a, \pi) : f(x) = \int_{S^d} a(w) \sigma(w^\top x) d\pi(w) \right\}$$

- ▶ Barron norm $\gamma_p(f)$:

$$\gamma_p(f) := \inf_{(a, \pi) \in \Theta_f} \left(\int_{S^d} |a(w)|^p d\pi(w) \right)^{1/p}$$

- ▶ **Barron space**

$$\mathcal{B}_p(\Omega) = \{f(x) : \gamma_p(f) < \infty\}$$

- ▶ Since π is a probability distribution, by Hölder's inequality

$$\gamma_p(f) \leq \gamma_q(f), \quad \text{if } q \geq p > 0.$$

and thus

$$\mathcal{B}_\infty(\Omega) \subset \cdots \subset \mathcal{B}_2(\Omega) \subset \mathcal{B}_1(\Omega)$$

New Results on GD Convergence of NNs

Relaxed the over-parametrization assumption

- $m = \Omega(n^6/(\delta^3\lambda_0^4))$ of Du et al. (2018) is relaxed to
 $m = \Omega(n^2/(\delta\lambda_0^4))$

Theorem 3.2: For any $0 < \delta < 1$, assume $m = \Omega(n^2 \ln(n^2/\delta)/(\delta\lambda_n^4))$, then, with probability at least $1 - \delta$, over the random initialization, we have, for all $t \geq 0$,

$$\hat{L}_n(t) \leq e^{-m\lambda_n t} \hat{L}_n(0),$$

where $\lambda_n > 0$ is the assumption on the lower bound of the corresponding Gram matrices.

Same as Fixed Kernel Method

- ▶ The over-parametrized GD training is the same as fixed kernel method

- ▶ **Definitions:**

- ▶ Let $\hat{f}_{ker}(x, t) := f_m(x; a_t, w_0),$

where only a_t is optimized and w_0 is left unchanged from its initial value.

- ▶ and let

$$\hat{f}(x, t) := f_m(x; a_t, w_t),$$

where both w and linear weights, a , are optimized.

Theorem 3.3: Assume $m = \Omega(n^2 \ln(n^2/\delta)/(\delta \lambda_n^4))$ for $0 < \delta < 1$ and $\beta \leq 1$. Then with probability at least $1 - 6\delta$ we have

$$|\hat{f}(x, t) - \hat{f}_{ker}(x, t)| \lesssim \frac{c_\delta^2}{m} \left(\frac{1}{\sqrt{m}} + \beta + \sqrt{m}\beta^3 \right)$$

where $c_\delta = 1 + \sqrt{\ln(1/\delta)}$

Remark If $\beta = o(m^{-1/6})$, then the right-hand side goes to zero

Relaxed Over-parametrization and Early Stopping

- ▶ Remove the over-parametrization assumption
- ▶ Assume that target function, f^* , is Barron's
- ▶ and $\|f^*\|_\infty \leq 1$

Theorem 4.1: Let $\beta = c/m$, and assume f^* is Barron function, with $\|f^*\|_\infty \leq 1$. then, for any $0 < \delta < 1$, with probability at least $1 - 4\delta$ we have

$$\hat{L}_n(t) \leq C \left(\frac{1}{m} + \frac{1}{mt} + \frac{1}{\sqrt{n}} \right)$$

Corollary (Generalization with early stopping) Assume $m > n$ and let $t = \sqrt{n}/m$. Then, under the assumption of Theorem 4.2,

$$L(t) \lesssim \left(\frac{1}{m} + \frac{1}{\sqrt{n}} \right)$$

Implicit Bias and Margin Maximization of Gradient Descent

- ▶ Could it be that we are finding nice generalizable solutions due to GD optimization?
- ▶ For linear predictors with linearly separable data, Soudry, Hoffer, and Srebro (2017) show that GD on the cross-entropy loss is implicitly biased towards a maximum margin direction.
 - ▶ Bias of GD towards margin maximization means that gradient descent "prefers" a solution which is likely to generalize well, and not just achieve low empirical risk.
- ▶ The preceding work inspired many other results, e.g.: Ji and Telgarsky 2019+; Gunasekar et al. 2018; Lyu and Li 2019; Chizat and Bach 2020; Ji et al. 2020.
- ▶ In this lecture, I'll cover some of the results from Chapter 10 in Telgarsky, 2021 monograph.

Implicit Margin Maximization of Gradient Descent

- ▶ Consider n data points: $(y_i, x_i), 1 \leq i \leq n, y_i \in \{-1, 1\}, x_i \in \mathbb{R}^d$
(For convenience, assume $x_i \equiv \tilde{x}_i = (1, x_{i1}, \dots, x_{i,d-1})$)
- ▶ Assume that data points are *linearly separable*, i.e., there exists w , such that

$$\min_i y_i \langle w, x_i \rangle > 0$$

- ▶ Typically, there are many such w -s. Which one is the best in terms of generalization?
- ▶ **Maximum margin classifier** is given by $\hat{y}(x) = \text{sign} \langle w^*, x \rangle$, where

$$w^* = \underset{\|w\|_2=1}{\operatorname{argmax}} \min_i y_i \langle w, x_i \rangle$$

and the **margin** is $\gamma := \min_i y_i \langle w^*, x_i \rangle$.

- ▶ What is the geometric interpretation of this classifier? Why do we expect it to generalize well? (This is the basis of Support Vector Machines.)
- ▶ It turns out that under the appropriate conditions, gradient descent converges to the maximum margin classifier.

Implicit Margin Maximization of Gradient Descent

- ▶ Consider now replacing $\langle w, x_i \rangle$ with $f(x_i; w)$, and **margin mapping**

$$m_i(w) := y_i f(x_i; w),$$

where $f(x)$ is locally-Lipschitz and L-homogeneous, i.e.,
 $f(cx) = c^L f(x), c \geq 0$.

- ▶ For $\ell(z) = e^{-z}$, let \mathcal{L} be unnormalized loss (no division by n)

$$\mathcal{L}(w) := \sum_{i=1}^n \ell(m_i(w)) = \sum_{i=1}^n \ell(y_i f(x_i; w)).$$

- ▶ Next, we say that data is **m-separable** if there is a w , such that $\min_i m_i(w) > 0$.
- ▶ Now, define the (general) **margin, maximum margin and smooth margin**, respectively as

$$\gamma(w) := \min_i m_i(w/\|w\|) = \frac{\min_i m_i(w)}{\|w\|^L}, \quad \bar{\gamma} := \max_{\|w\|=1} \gamma(w), \quad \tilde{\gamma} := \frac{\ell^{-1}(\mathcal{L}(w))}{\|w\|^L}.$$

Motivation for smooth margin comes from (recall ℓ is exponential):
 $\ell^{-1}(\mathcal{L}(w)/n) \geq \min_i m_i(w) \geq \ell^{-1}(\max_i \ell(m_i(w))) \geq \ell^{-1}(\mathcal{L}(w))$

Implicit Margin Maximization of Gradient Descent

Note: Gradient descent is biased towards larger margins, which guarantee good generalization.

Theorem (10.1 in Telgarsky, 2021) Consider the linear case, with linearly separable data, exponential loss, $\ell(z) = e^{-z}$, and $\max_i \|x_i\| \leq 1$. Then, for gradient descent path w_t with $w(0) = 0$,

$$\gamma(w_t) \geq \tilde{\gamma}(w_t) \geq \bar{\gamma} - \frac{\ln n}{\ln t + \ln(2n\gamma^2) - \ln \ln(2tne\gamma^2)}$$

Proof: Consider gradient flow path which satisfy

$$\dot{w}(t) = -\nabla \mathcal{L}(w(t)),$$

which, for $u(t) := \ell^{-1}(\mathcal{L}(w(t)))$, imply ($\ell' = -\ell$ and $\ell^{-1}(z) = -\ln z$)

$$\dot{u}(t) = \left\langle \frac{-\nabla \mathcal{L}(w(t))}{\mathcal{L}(w(t))}, \dot{w}(t) \right\rangle = \frac{\|\dot{w}(t)\|^2}{\mathcal{L}(w(t))}.$$

Implicit Margin Maximization of Gradient Descent

Proof: Now we can lower bound $\gamma(w(t))$ as

$$\gamma(w(t)) \geq \tilde{\gamma}(w(t)) = \frac{u(t)}{v(t)} = \frac{u(0)}{v(t)} + \frac{\int_0^t \dot{u}(s) ds}{v(t)}, \quad (1)$$

where $v(t) := \|w(t)\|$ and $u(t)$ was previously defined. Now, we bound the second term in the preceding equation. To this end, note

$$\begin{aligned} \|\dot{w}(s)\| &\geq \langle \dot{w}(s), w^* \rangle = \left\langle -\sum_i x_i y_i \ell'(m_i(w(s))), w^* \right\rangle \\ &= \sum_i \ell(m_i(w(s))) \langle x_i y_i, w^* \rangle, \quad (\ell' = -\ell) \\ &\geq \gamma \mathcal{L}(w(s)) \end{aligned}$$

Also,

$$v(t) = \left\| \int_0^t \dot{w}(s) ds \right\| \leq \int_0^t \|\dot{w}(s)\| ds$$

Implicit Margin Maximization of Gradient Descent

Proof: Combining the preceding, we lower bound the second term in (1)

$$\frac{\int_0^t \dot{u}(s) ds}{v(t)} = \frac{\int_0^t \frac{\|\dot{w}(s)\|^2}{\mathcal{L}(w(s))} ds}{v(t)} \geq \gamma \frac{\int_0^t \|\dot{w}(s)\| ds}{v(t)} \geq \gamma$$

Since the above inequality holds for any margin γ , it holds in for the maximum margin $\bar{\gamma}$.

Next, for the first term, $u(0)/v(t)$, in (1), note that $\mathcal{L}(w(0)) = n$, and thus $u(0) = -\ln n$, and it remains to prove that

$$\|w(t)\| \geq \ln(t) + \ln(2n\gamma^2) - 2 \ln \ln(2tne\gamma^2).$$

In this regard, we first prove the following lemma.

Lemma For any convex loss function \mathcal{L} and any $z \in \mathbb{R}^d$, along the gradient flow path $w(t)$, we have

$$\mathcal{L}(w(t)) \leq \mathcal{L}(z) + \frac{1}{2t} (\|w(0) - z\|_2^2 - \|w(t) - z\|_2^2).$$

Implicit Margin Maximization of Gradient Descent

Proof of the lemma: note that

$$\begin{aligned}\frac{1}{2}(\|w(t) - z\|_2^2 - \|w(0) - z\|_2^2) &= \frac{1}{2} \int_0^t \frac{d}{ds} \|w(s) - z\|_2^2 ds \\ &= \int_0^t \left\langle \frac{dw}{ds}, w(s) - z \right\rangle ds \quad \left(\frac{dw}{ds} = -\nabla \mathcal{L}(w(s)) \right) \\ &= \int_0^t \langle -\nabla \mathcal{L}(w(s)), w(s) - z \rangle ds \\ &\leq \int_0^t (\mathcal{L}(z) - \mathcal{L}(w(s))) ds \quad (\text{convexity}) \\ &\leq t\mathcal{L}(z) - t\mathcal{L}(w(t)),\end{aligned}$$

where in the second to the last inequality we used the gradient flow assumption $dw(t)/dt = -\nabla \mathcal{L}(w(t))$, and the convexity assumption

$$\mathcal{L}(z) \geq \mathcal{L}(w) + \langle \nabla \mathcal{L}(w), z - w \rangle,$$

which concludes the proof of the lemma.

Implicit Margin Maximization of Gradient Descent

Proof: Now, we complete the proof of the theorem, by applying the preceding lemma with $z = \ln(2tn\gamma^2)w^*/\gamma$, and recall $w(0) = 0$

$$\begin{aligned} n\ell(\|w(t)\|) &\leq n \min_i \ell(y_i \langle w(t), x_i \rangle) \leq \mathcal{L}(w(t)) \\ &\leq \mathcal{L}(z) + \frac{1}{2t} (\|w(0) - z\|_2^2 - \|w(t) - z\|_2^2) \\ &\leq \mathcal{L}(z) + \frac{\|z\|_2^2}{2t} \\ &\leq \frac{n}{2tn\gamma^2} + \frac{(\ln(2tn\gamma^2))}{2t\gamma^2}, \end{aligned}$$

which, after dividing by n , taking $\ell^{-1}(z) = -\ln z$ on both sides, implies

$$\|w(t)\| \geq \ln(2tn\gamma^2) - \ln(1 + \ln(2tn\gamma^2)) = \ln(t) + \ln(2n\gamma^2) - \ln \ln(2tne\gamma^2),$$

which completes the proof of the theorem.

Implicit Margin Maximization of Gradient Descent


Few remarks:

- ▶ The rate of convergence is $1/\ln(t)$, which can be accelerated by rescaling time, Chizat and Bach (2020): see Theorem 10.2 in Telgarsky, 2021.
- ▶ Extension to nonlinear homogeneous functions can be found Theorem 10.3 in Telgarsky, 2021. The theorem is originally due to Lyu and Li (2019).
- ▶ Extensions to NNs under restrictions can be found in the aforementioned papers:
 - ▶ [The Implicit Bias of Gradient Descent on Separable Data](#), Soudry et al., 2017.
 - ▶ [Gradient Descent Maximizes the Margin of Homogeneous Neural Networks](#), Lyu, Kaifeng, and Jian Li, 2019.
 - ▶ [Implicit Bias of Gradient Descent for Wide Two-Layer Neural Networks Trained with the Logistic Loss](#), Chizat and bach, 2020.
 - ▶ For additional references check the follow up references on Google Scholar, and Chapter 10 in Telgarsky 2021.

Deep Residual Networks

Proposed by: [Deep Residual Learning for Image Recognition](#), He et al., 2015.

- ▶ Easier to optimize at larger depth: scale to 100+ depth
- ▶ Reduces the problem of vanishing/exploding gradients
- ▶ basic building block



Deep Residual Networks

- ▶ Hidden state transformations

$$\mathbf{h}_{t+1} = \mathbf{h}_t + f(\mathbf{h}_t, \theta_t), \quad (2)$$

$$t \in \{0, \dots, T\}, \mathbf{h}_t \in \mathbb{R}^d$$

- ▶ \mathbf{h}_0 is the input layer and \mathbf{h}_T is the output layer

$$\mathbf{h}_T = \mathbf{h}_0 + \sum_{t=0}^{T-1} f(\mathbf{h}_t, \theta_t)$$

The function is additive and the gradient should behave better

Continuous Approximation to ResNets

Neural Ordinary Differential Equations, by Chen et al., 2018.

- Approximate Equation (2) by a differential equation

$$\frac{d\mathbf{h}(t)}{dt} = f(\mathbf{h}(t), t, \theta)$$




Figure 1: *Left:* A Residual network defines a discrete sequence of finite transformations. *Right:* A ODE network defines a vector field, which continuously transforms the state. *Both:* Circles represent evaluation locations.

Continuous Approximation

Claimed advantages

- ▶ Can use standard ODE solvers
- ▶ Memory efficiency
- ▶ Solving ODEs well understood: over 120 years of experience
- ▶ Parameter efficiency
- ▶ Continuous transformations and change of variables
- ▶ Continuous time series models: can incorporate data that arrives at arbitrary times

Computing the Gradient

- ▶ Instead of backpropagation, compute the gradient by solving another (adjoint) ODE backward in time
- ▶ Consider optimizing a scalar loss $L(\cdot)$, whose input is the result of an ODE solver

$$L(z(t_1)) = L\left(z(t_0) + \int_{t_0}^{t_1} f(z(t), t, \theta) dt\right)$$

- ▶ To optimize L , we need a gradient with respect to θ
- ▶ So, we first compute the *adjoint*

$$\mathbf{a}(t) = \frac{\partial L}{\partial z(t)}$$

- ▶ $\mathbf{a}(t)$ dynamics is given by another ODE (an instantaneous analog of the chain rule)

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^\top \frac{\partial f(z(t), t, \theta)}{\partial z(t)}$$

Computing the Gradient

- ▶ $\mathbf{a}(t)$ dynamics is given by another ODE (an instantaneous analog of the chain rule)

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^\top \frac{\partial f(z(t), t, \theta)}{\partial z(t)}$$

- ▶ We can compute $\mathbf{a}(t) = \partial L / \partial z(t)$ by solving the preceding equation backward in time starting with the initial value $\mathbf{a}(t_1) = \partial L / \partial z(t_1)$
- ▶ Complication: need to know $z(t)$ along its entire trajectory
 - ▶ Recompute $z(t)$ backward in time, together with the adjoint
- ▶ Finally, compute the gradient by evaluating the integral

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial f(z(t), t, \theta)}{\partial z(t)} dt$$

Reading

- ▶ PAC Learning and Generalization Theory - [ML2014] book:
VCdim: Chapters 6&28; VCdim of NNs: Theorem 20.6 in Ch. 20.
- ▶ Generalization bounds
 - ▶ A Comparative Analysis of the Optimization and Generalization Property of Two-layer Neural Network and Random Feature Models Under GD, Weinan et al., Feb, 2020.
- ▶ Implicit bias of gradient descent
 - ▶ Chapter 15 in Telgarsky, 2021.
 - ▶ The Implicit Bias of Gradient Descent on Separable Data, Soudry et al., 2017.
 - ▶ Gradient Descent Maximizes the Margin of Homogeneous Neural Networks, Lyu, Kaifeng, and Jian Li, 2019.
 - ▶ Implicit Bias of Gradient Descent for Wide Two-Layer Neural Networks Trained with the Logistic Loss, Chizat and Bach, 2020.
 - ▶ For additional references check the follow up references on Google Scholar, and Chapter 15 in Telgarsky 2021.
- ▶ Residual Networks
 - ▶ Deep Residual Learning for Image Recognition, He et al., 2015.
 - ▶ Neural Ordinary Differential Equations, by Chen et al., 2018.

Final Project

The key difference from other courses and guiding questions:

- ▶ What did you learn about a neural network?
 - ▶ The focus should be on NN properties instead of applications.
- ▶ How do the changes in NN impact its performance?
 - ▶ The changes could be: architecture (e.g., width/depth), activation function, training method, normalization, dropout...
 - ▶ In class, we focused on plain vanilla feed-forward networks, but you could choose other types, e.g., ResNets.
- ▶ You could center your questions on one or more of the general themes we focused on in class:
 1. Approximation and interpolation theory and the impact of depth.
 2. Optimization landscape and global convergence.
 3. Generalization theory: conditions for small/bounded testing errors.
- ▶ Many of the problems we formulated in the context of wide/over-parametrized networks with two types of scaling: NTK/lazy training or mean-field/active training.

Final Project

Rough Paper Outline, about 15 pages:

1. **Introduction:** e.g., describe the general problem area, DL and specific subtopic(s). Brief literature review, etc.
2. **Detailed Problem Description:** More detailed literature review for a selected problem(s), detailed description of the known results, theoretical or experimental, etc.
3. **Some Reproduction:** Theoretical or Experimental partial or full reproduction of the results. For example, run some simulations that illustrate main results.
4. **New Results: Theoretical or Experimental** Describe in detail your results. If experimental, describe the experiments and results. Explain clearly the graphs and tables from experiments, etc.
5. **Discussion and conclusion:** e.g., try to draw general inferences from your results. Compare to the known results from the literature, etc.

Final Project

- ▶ **Deliverables:**

- ▶ **Paper:** about 15 pages - the most important part.
- ▶ **Presentations:** about 10min each, 10 slides
- ▶ **Software:** Document your code well

- ▶ **First set presentations:** April 25, during the last class
3% EC for those presenting on April 25
- ▶ Additional presentation slots during study/exam week: TBA
- ▶ **Project due:** During the exam week of May 5-12: TBA
- ▶ Academic Honesty - do not plagiarize; Turnitin will be used to check for originality

Have Fun and GOOD LUCK!